



Getting Started with HBase Application development

Sridhar Reddy

sreddy@mapr.com

March 17th, 2015



Objectives of this session

- What is HBase?
 - Why do we need NoSQL / HBase?
 - Overview of HBase & HBase data model
 - HBase Architecture and data flow
 - Hbase Use Cases
- How to get started
 - Demo/Lab using HBase Shell using MapR Sandbox
- Developing Applications using HBase Java API
 - Demo/Lab to perform CRUD operations using put, get, scan, delete



Agenda

- Why do we need NoSQL / HBase?
- Overview of HBase & HBase data model
- HBase Architecture and data flow
- HBase Use Cases
- Demo/Lab using HBase Shell
 - Create tables and CRUD operations using MapR Sandbox
- **HBase Java API to perform CRUD operations**
 - Demo HBase Java API using Eclipse & MapR Sandbox



Prerequisite for Hands-On-Labs

- Install a one-node MapR Sandbox on your laptop
- Install and configure Eclipse to develop HBase applications using Java API
- MapR Client is optional

Hbase_Tutorial_3_2015.pdf

GettingStartedWithHBase_3_17_2015.pdf

exercises.zip &

solutions.zip



Why do we need NoSQL / HBase?

Relational Database is **typed** and **structured before stored**:

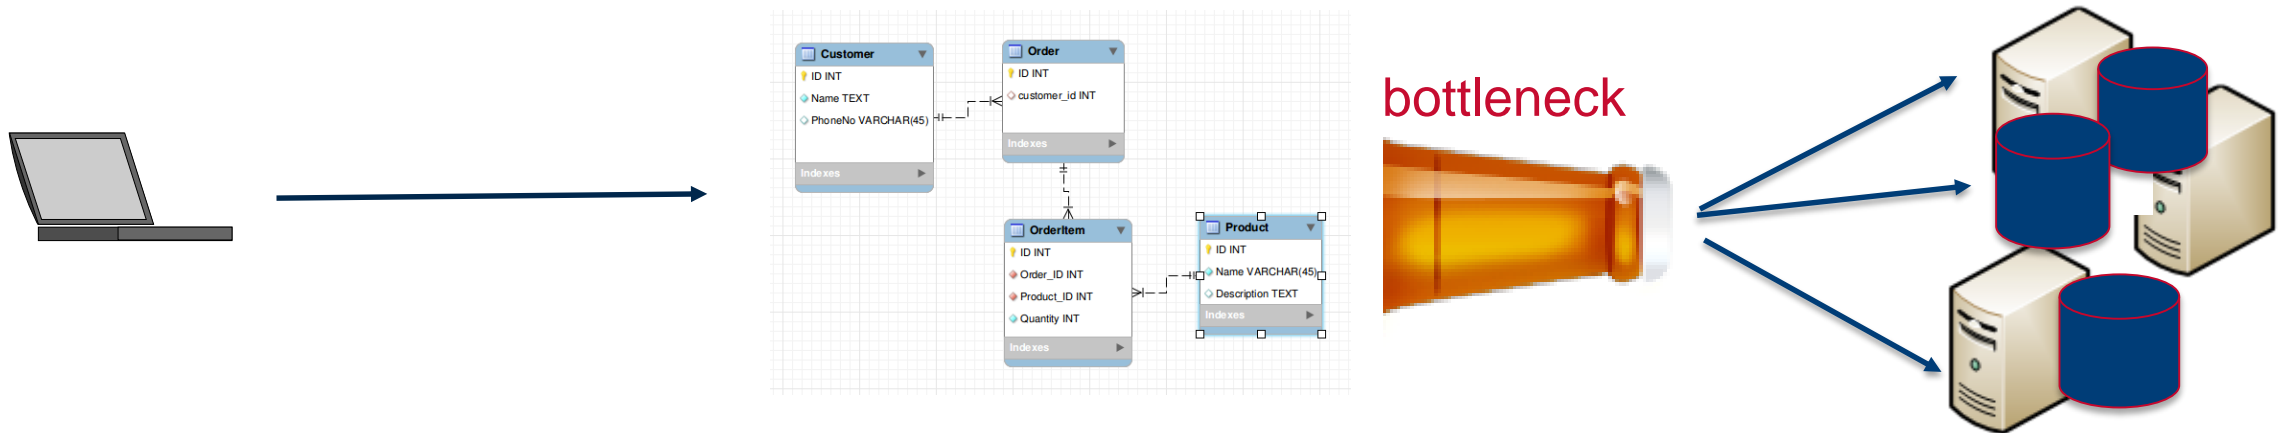
- Entities map to tables, normalized
- **Transactions** handle concurrency , consistency
- Structured Query Language
 - Joins tables to bring back data

Customer Table							Customer-Order Junction Table		Order Table				
Customer							Customer Order		Order				
ID	First Name	Last Name	Address	City	State	Zip Code	ID	Number	Number	Item	Quantity	Unit Price	Total Price
1	John	Doe	1 Junction Avenue	San Jose	CA	95134	1	1001	1001	Hamilton Beach Coffee Maker	1	\$4.99	\$54.99
2	Albert	Einstein	2 Physics Place	New York	NY	10003	1	1002	1001	Kindle Fire 7"	1	\$139.00	\$139.00
3	Ben	Franklin	1776 Freedom Way	Philadelphia	PA	19106	2	1003	1001	Divergent (paperback)	1	\$5.49	\$5.49
4	Mark	Twain	47 Huckleberry Drive	Hannibal	MO	63401	3	1004	1001	Seagate Desktop HDD 4TB	1	\$156.95	\$156.95
5	Abe	Lincoln	16 Springfield Road	Springfield	IL	62561	4	1005	1002	In Paradise: A Novel	1	\$21.49	\$21.49
							5	1006	1003	Astonish Me: A Novel	1	\$16.41	\$16.41
							5	1007	1003	Updike	1	\$21.77	\$21.77
									1003	Sous Chef: 24 Hours on the Line	1	\$18.85	\$18.85
									1003	Secrecy (paperback)	1	\$12.71	\$12.71
									1004	Happy (from Despicable Me 2)	1	\$1.29	\$1.29
									1004	Let It Go	1	\$1.29	\$1.29
									1005	Dark Horse	1	\$1.29	\$1.29
									1005	All of Me (Album Version)	1	\$1.29	\$1.29
									1005	The Man	1	\$1.29	\$1.29
									1005	Frozen (Two-Disc Blu-ray)	1	\$19.96	\$19.96
									1007	Anchorman 2: The Legend Continues	1	\$16.96	\$16.96



What changed to bring on NoSQL? Big data

Horizontal scale : partition or shard tables across cluster



Distributed Joins, Transactions are Expensive

- **Cons of the Relational Model:**
 - Does not **scale** horizontally:
 - Sharding is **difficult** to manage
 - Distributed join, transactions **do not scale** across shards

NoSQL Landscape

Key Value Store

- Couchbase
- Riak
- Citrusleaf
- Redis
- BerkeleyDB
- Membrain
- ...

Document

- MongoDB
- CouchDB
- RavenDB
- Couchbase
- ...

Graph

- OrientDB
- DEX
- Neo4j
- GraphBase
- ...

Wide Column

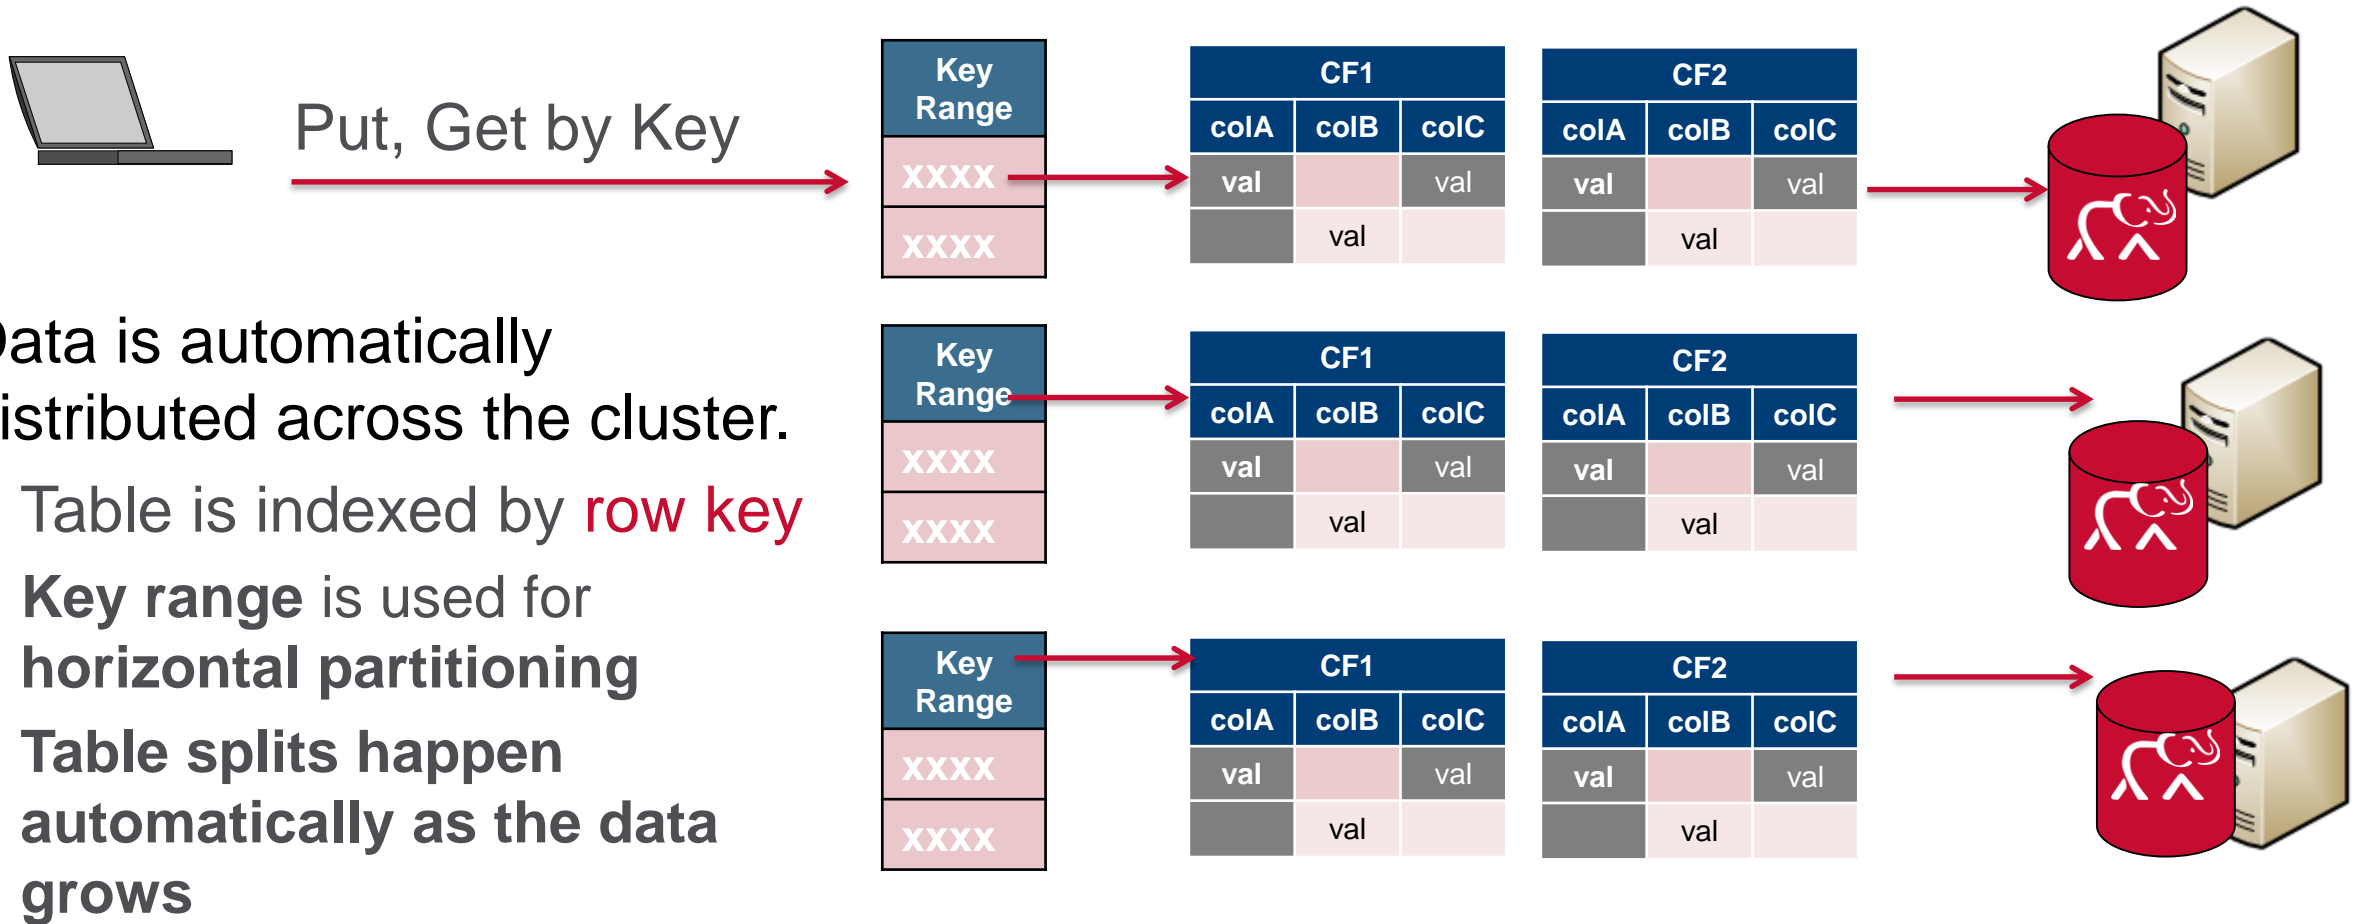
- **HBase**
- **MapR-DB**
- Hypertable
- Cassandra
- ...



Hbase designed for Distribution, Scale, Speed



HBase is a Distributed Database



HBase is a ColumnFamily oriented Database

Customer id		Customer Address data			Customer order data		
		CF1			CF2		
RowKey		colA	colB	colC	colA	colB	colC
axxx		Val		val	val		val
gxxx			val			val	

Data is **accessed and stored together by RowKey**


Similar Data is grouped & stored in Column Families

- share common properties:
 - Number of versions
 - Time to Live (TTL)
 - Compression [lz4, lzf, Zlib]
 - In memory option ...



HBase designed for Distribution

- Distributed data stored and accessed together:
 - **Key range** is used for **horizontal partitioning**
- Pros
 - **scalable** handles data volume and velocity
 - **Fast Writes and Reads by Key**
- Cons
 - **No** joins
 - Need to **know how** data will be **queried in advance** to do good schema design to achieve best performance



HBase Data Model

Row Keys: identify the rows in an HBase table

Columns are grouped into column families

	Row Key	CF1			CF2				...
		colA	colB	colC	colA	colB	colC	colD	
R1	axxx	val		val	val			val	
	...								
	gxxx	val			val	val	val		
R2	hxxx	val	val	val	val	val	val	val	
	...								
	jxxx	val							
R3	kxxx	val		val	val			val	
	...								
	rxxx	val	val	val	val	val	val		



HBase Data Storage - Cells

- Data is stored in **Key value** format
- Value for each **cell** is specified by complete coordinates:
- (Row key, ColumnFamily, Column Qualifier, timestamp) => Value
 - RowKey:CF:Col:Version:Value
 - smithj:data:city:1391813876369:nashville

Cell Coordinates= Key				Value
Column Key				
Row key	Column Family	Column Qualifier	Timestamp	Value
Smithj	data	city	1391813876369	nashville



Logical Data Model vs Physical Data Storage

Logical Model

RowKey	CF1		CF2	
	colA	colB	colA	colC
ra	1		2	
rxxxx				
rxxx				

- Data is stored in Key Value format
- **Key Value** is stored for each **Cell**
- Column families data are stored in separate files

Key			Value
Row Key	CF1:Col	version	value
ra	cf1:cola	1	1

Physical Storage

Key			Value
row Key	CF2:Col	version	value
ra	cf2:cola	1	2

Physical Storage



Sparse Data with Cell Versions

	CF1:colA	CF1:colB	CF1:colC
Row1	<div>@time7: value3</div> <div>@time1: value1</div>		
Row10	<div>@time2: value1</div>	<div>@time2: value1</div>	
Row11	<div>@time6: value2</div> <div>@time3: value1</div>		
Row2	<div>@time4: value1</div>		<div>@time4: value1</div>



Versioned Data

- Version
 - **each** put, delete adds new cell, new version
 - A **long**
 - by **default** the current **time** in milliseconds if no version specified
 - **Last 3** versions are stored by **default**
 - **Configurable** by column family
 - You can **delete** specific cell **versions**
 - When a cell exceeds the **maximum** number of versions, the extra records are removed

Key	CF1:Col	version	value
ra	cf1:cola	3	3
ra	cf1:cola	2	2
ra	cf1:cola	1	1



Table Physical View

Physically data is stored per Column family as a sorted map

- Ordered by **row key, column qualifier** in **ascending** order
- Ordered by **timestamp** in **descending** order

Sorted by
Row key
and Column

Row key	Column qualifier	Cell value	Timestamp (long)
Row1	CF1:colA	value3	<i>time7</i>
Row1	CF1:colA	value2	<i>time5</i>
Row1	CF1:colA	value1	<i>time1</i>
Row10	CF1:colA	value1	<i>time4</i>
Row 10	CF1:colB	value1	<i>time4</i>

Sorted in
descending
order



Logical Data Model vs Physical Data Storage

Row Key	CF1		CF2	
	ca	cb	ca	cd
ra	1		2	
rb		3,4		
rc	5		6,7	8

Logical Model

Column families are stored separately
Row keys, Qualifiers are sorted lexicographically

Physical Storage

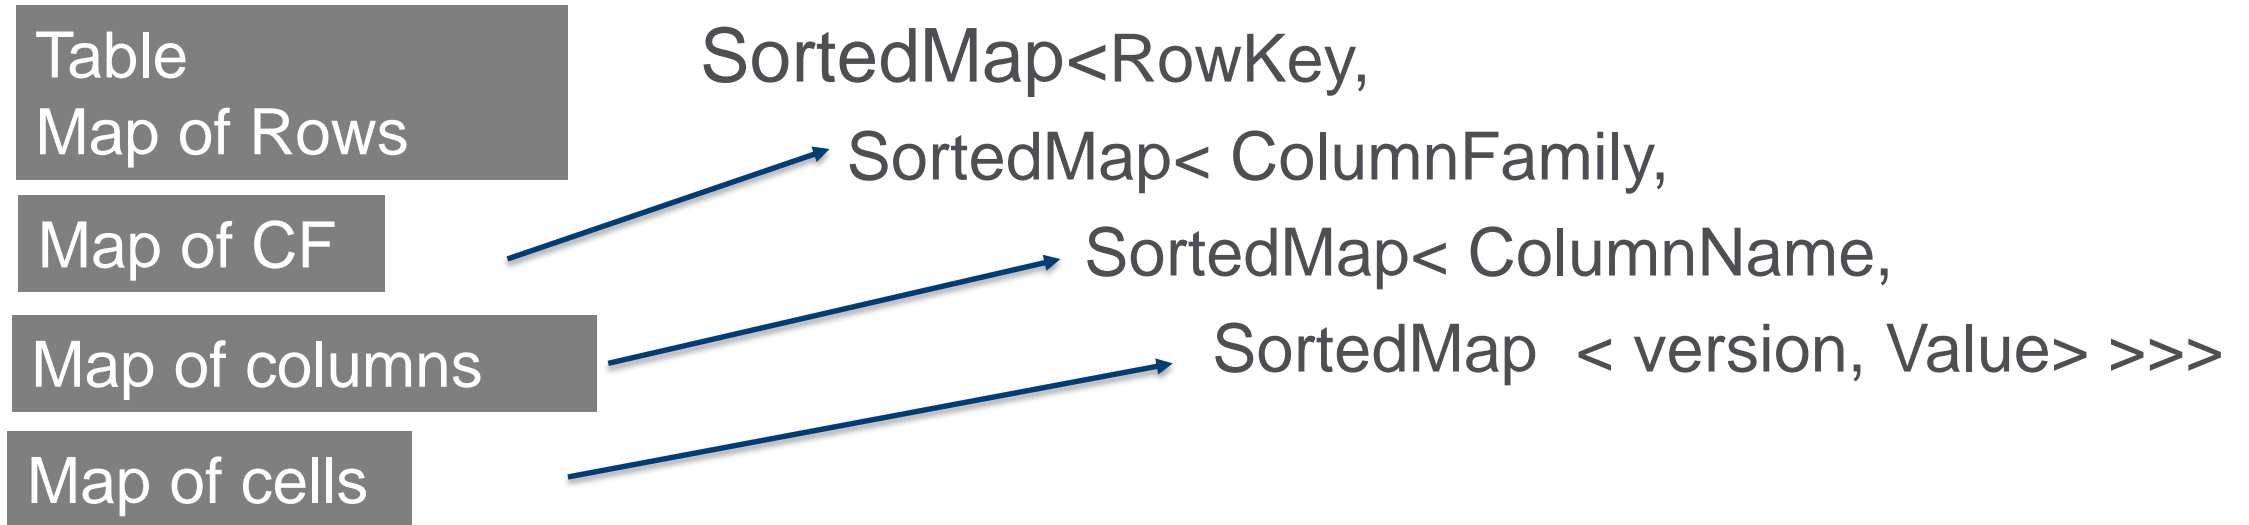
Key		Value	
Key	CF1:Col	version	value
ra	cf1:ca	1	1
rb	cf1:cb	2	4
rb	cf1:cb	1	3
rc	cf1:ca	1	5

Key		Value	
Key	CF2:Col	version	value
ra	cf2:ca	1	2
rc	cf2:ca	2	7
rc	cf2:ca	1	6
rc	cf2:cd	1	8



HBase Table is a Sorted map of maps

SortedMap<Key, Value>



Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8



HBase Table SortedMap<Key, Value>

<ra,<cf1, <ca, <v1, 1>>
 <cf2, <ca, <v1, 2>>>
<rb,<cf1, <cb, <v2, 4>
 <v1, 3>>>
<rc,<cf1, <ca, <v1, 5>>
 <cf2, <ca, <v2, 7>>
 <ca, <v1, 6>>
 <cd, <v1, 8>>>

Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8



Basic Table Operations

- **Create Table, define Column Families before data is imported**
 - but not the rows keys or number/names of columns
- Low level API, technically more demanding
- **Basic data access operations (CRUD):**
 - put Inserts data into rows (both create and update)
 - get Accesses data from one row
 - scan Accesses data from a range of rows
 - delete Delete a row or a range of rows or columns



HBase Architecture

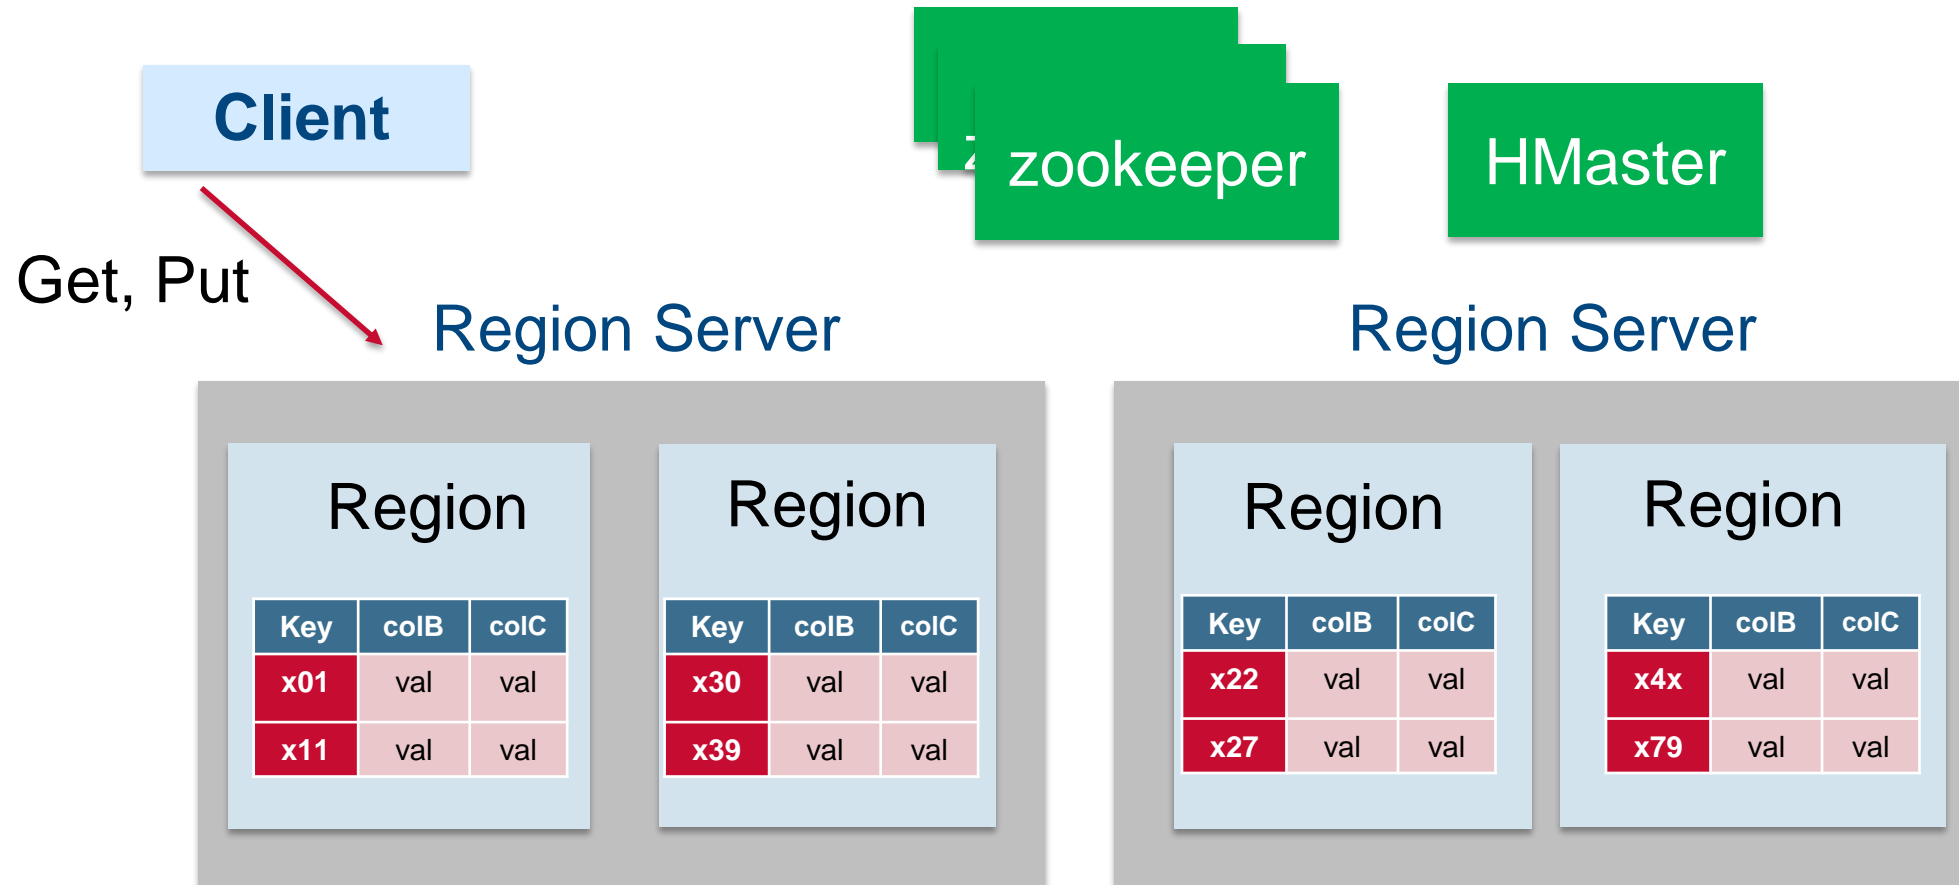
Data flow for Writes, Reads

Designed to Scale




What is a Region?

- Tables are partitioned into **key ranges (regions)**
- Region servers serve data for reads and writes
 - For the **range of keys** it is responsible for



Region Server Components

- WAL: write ahead log on **disk** (commit log), Used for recovery
- BlockCache: **Read** Cache, **L**east **R**ecently **U**sed evicted
- MemStore: **Write** Cache, sorted keyvalue updates.
- Hfile: sorted KeyValues on **disk**



HBase Write Steps

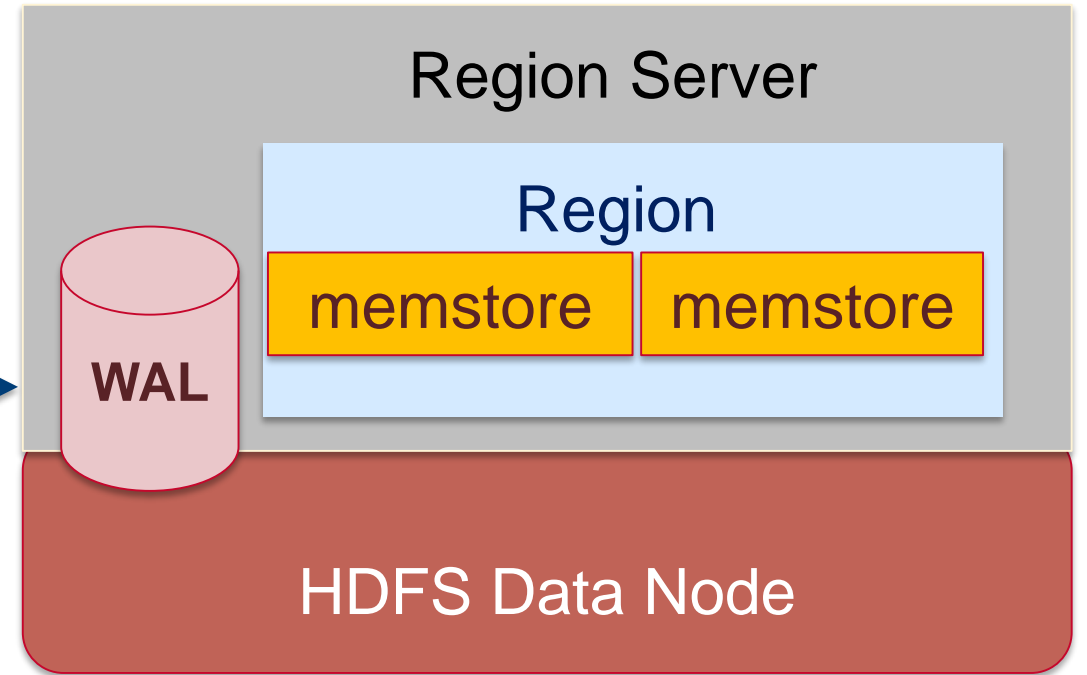


Put



each incoming record written to **WAL** for **durability**:

- log on disk
- updates appended sequentially



HBase Write Steps



Put

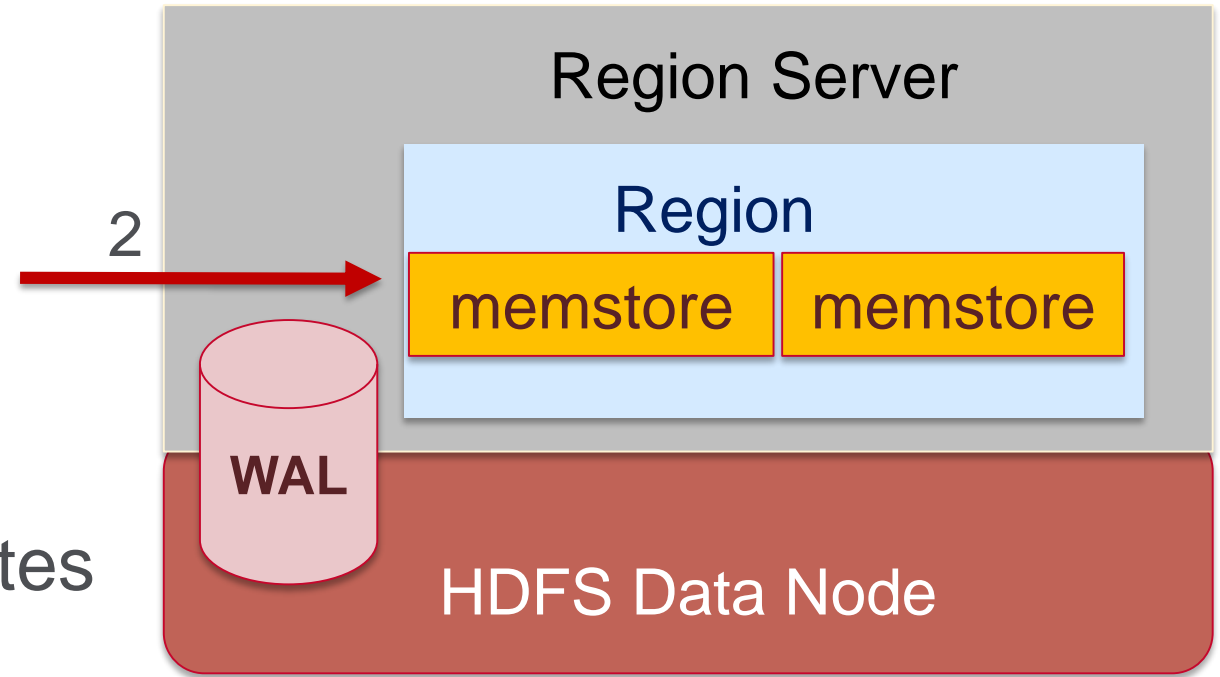


Next updates are written to the **Memstore**:

- write cache
- in-memory
- **sorted list of KeyValue** updates



Ack



Updates **quickly sorted in memory** are available to queries after put returns

HBase Memstore

- **in-memory**
- **sorted** list of **Key** → **Value**
- One per column family
- Updates **quickly sorted in memory**

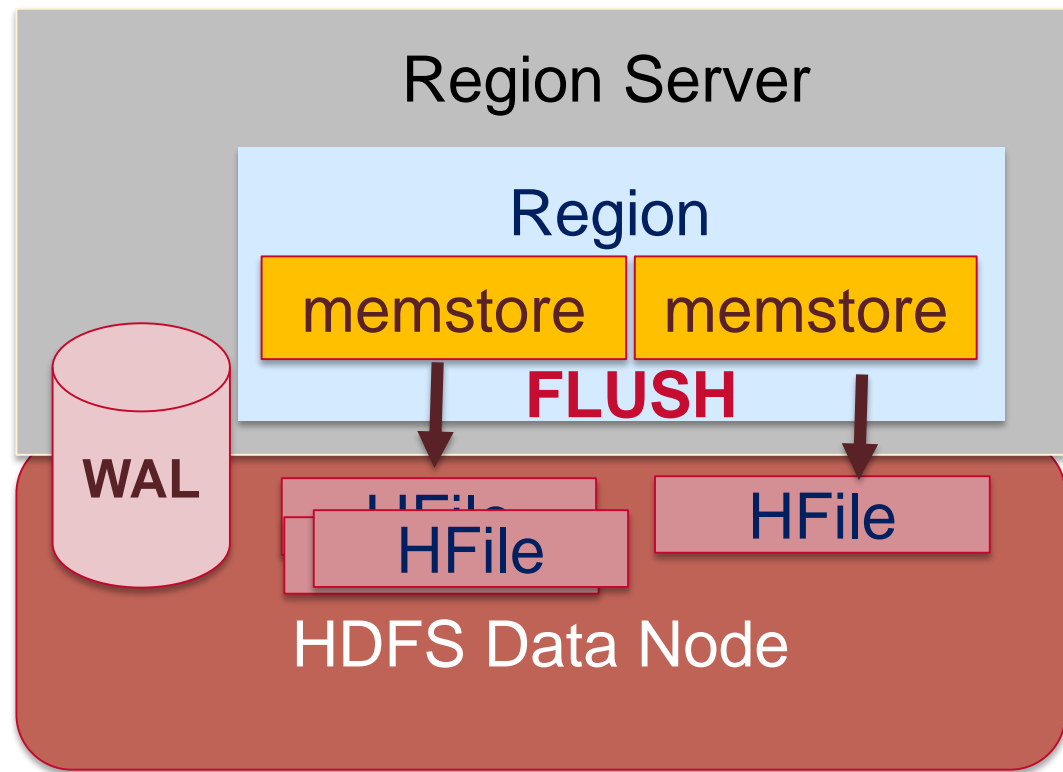


Key			Value
Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Key			Value
Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8



HBase Region Flush





When 1 **Memstore** is **full**:

- **all memstores** in region **flushed** to new **Hfiles** on **disk**
- Hfile: sorted list of key → values
On disk

HBase HFile

- On disk **sorted** list of key → values
- One per column family
- Flushed quickly to file
 - **Sequential write**

Sequential write

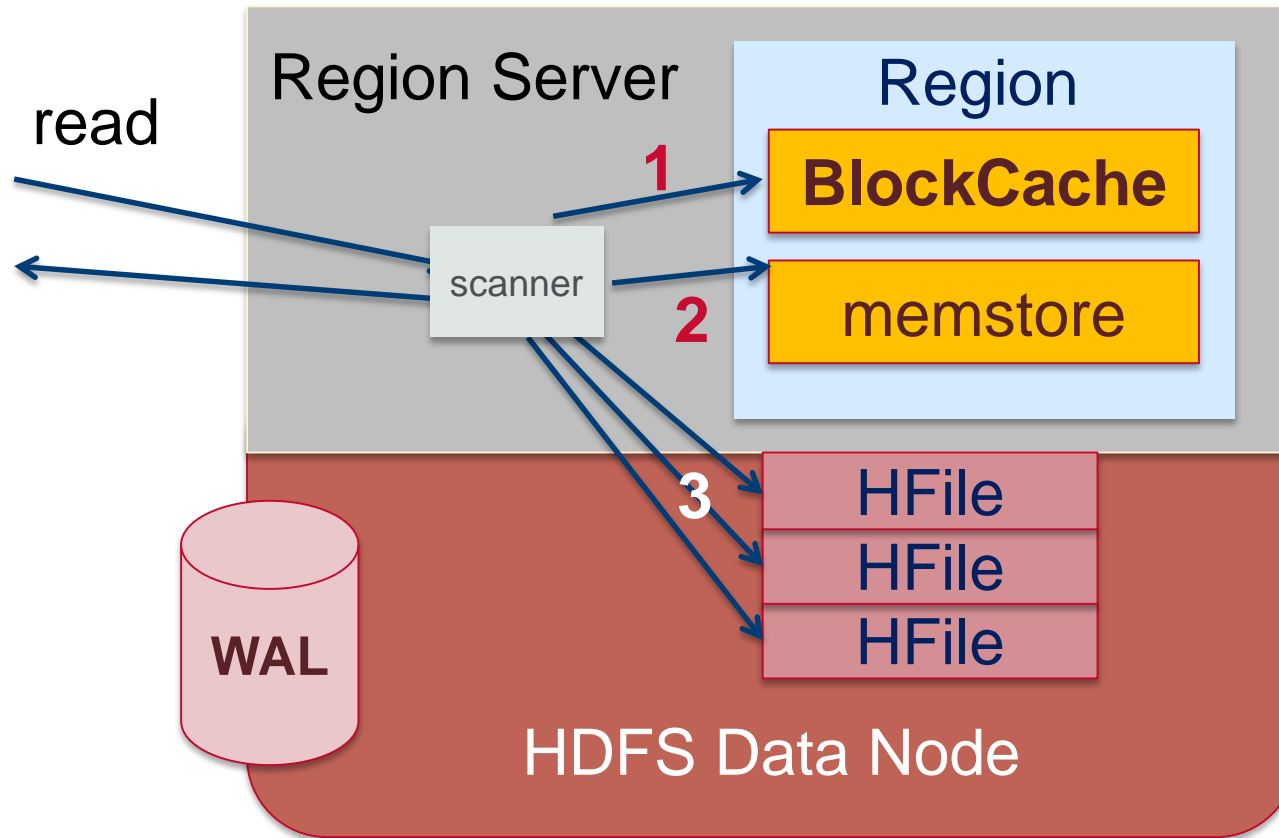


Key			Value
Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Key			Value
Key	CF2:Col	version	value
ra	cf2:ca	v1	2
rc	cf2:ca	v2	7
rc	cf2:ca	v1	6
rc	cf2:cd	v1	8



HBase Read Merge from Memory and Files

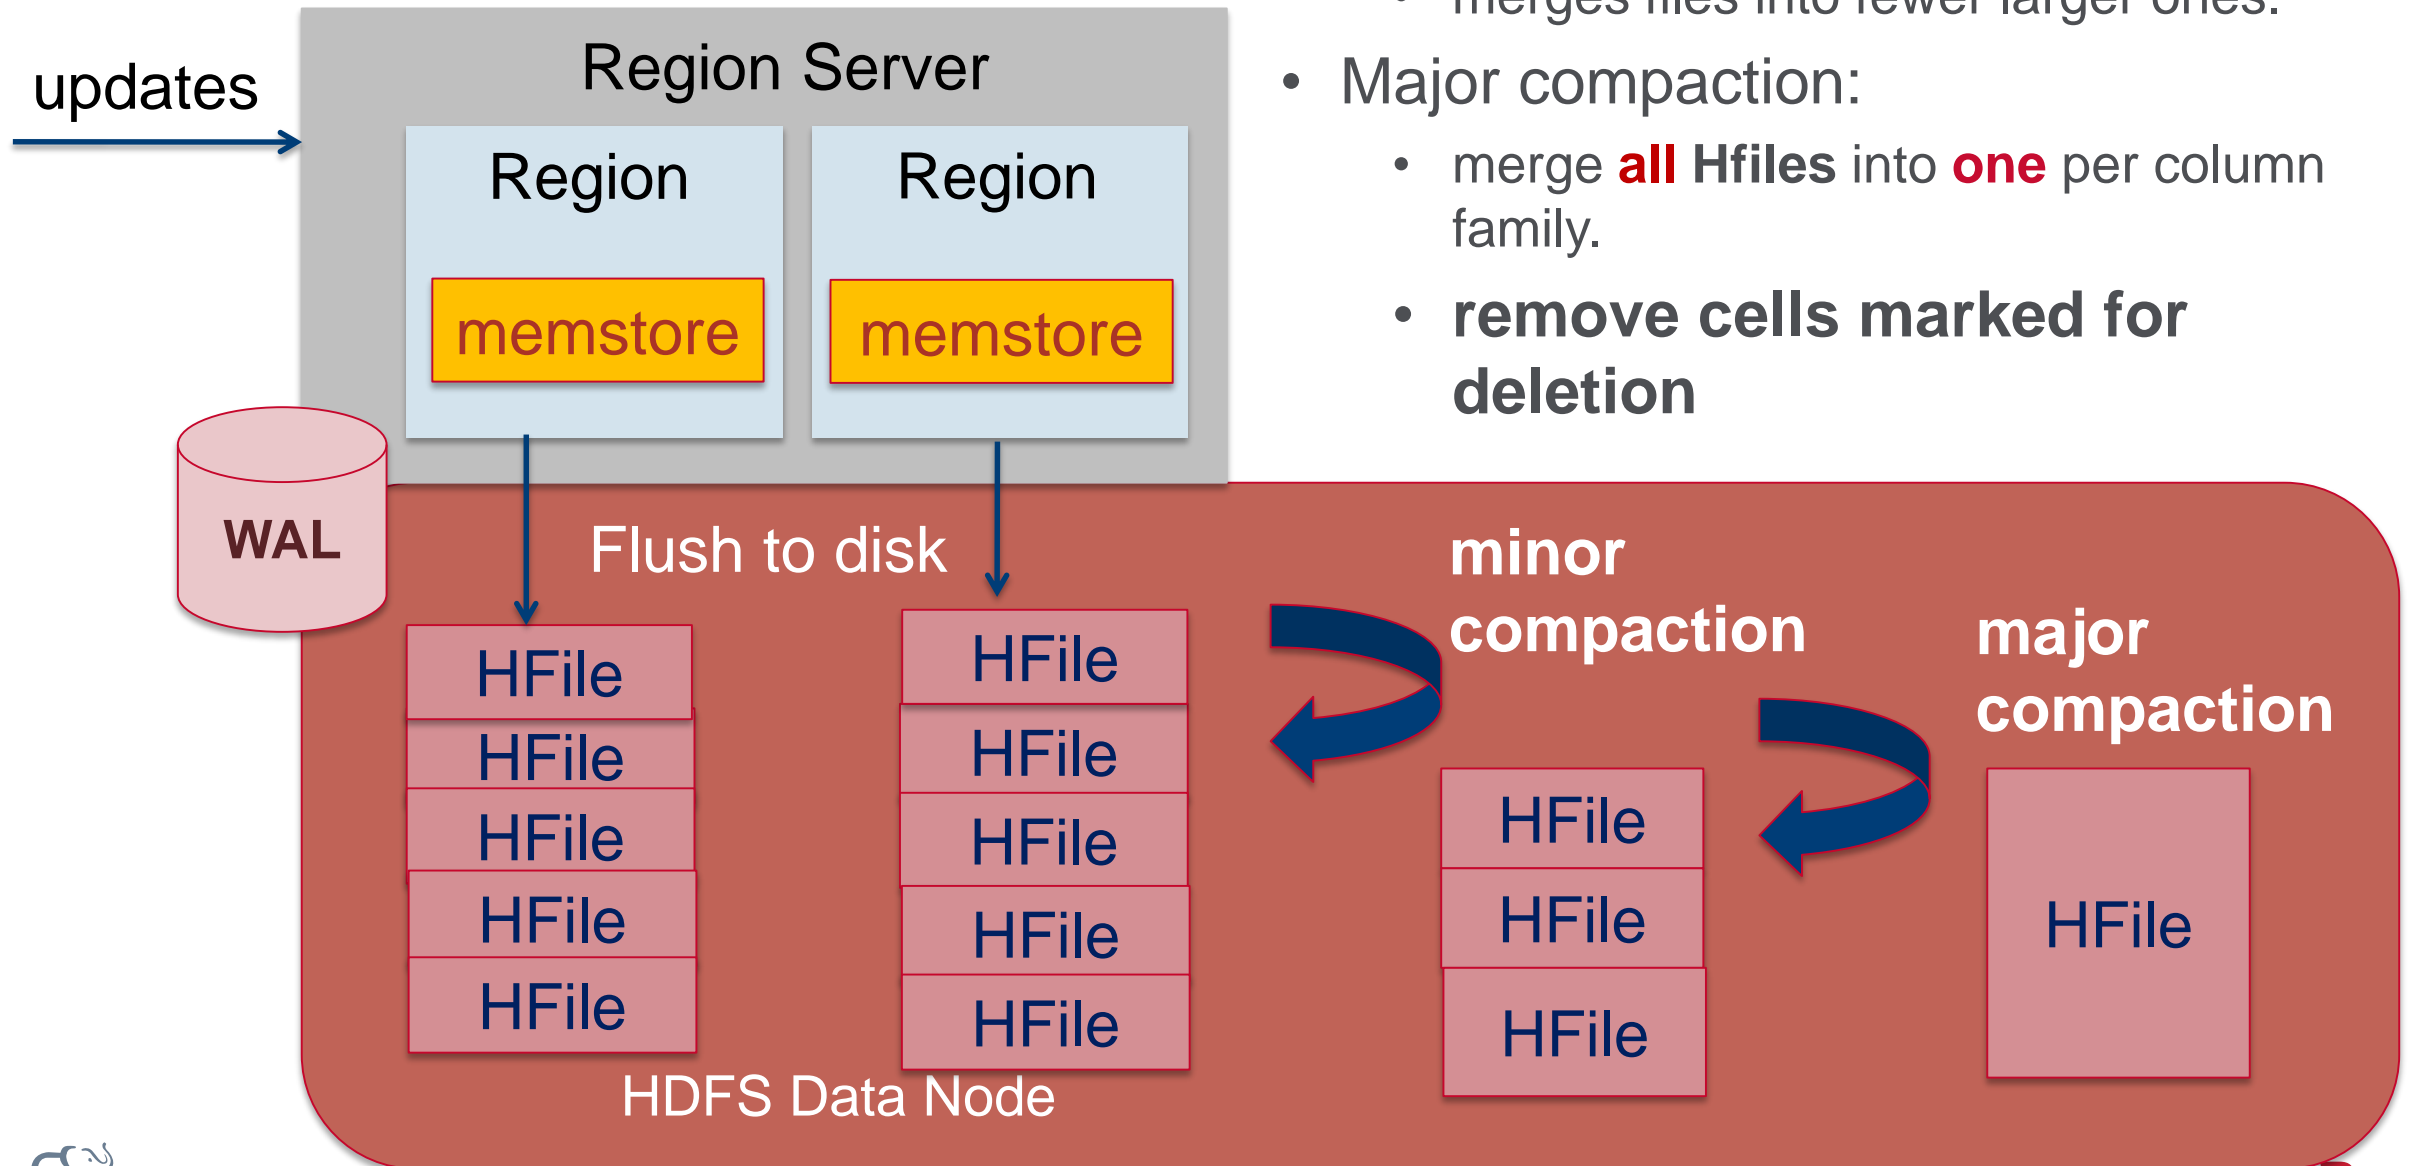


Get or Scan searches for Row Cell KeyValues:

1. Block Cache ((Memory)
2. Memstore (Memory)
3. Load HFiles from Disk into Block Cache based on indexes and bloomfilters


- MemStore creates multiple **small store files** over time when **flushing**.
- When a get/scan comes in, multiple files have to be examined

HBase Compaction




HBase Background: Log-Structured Merge Trees

- Traditional Databases use B+ trees:
 - **expensive** to update
- HBase: Log Structured Merge Trees
 - **Sequential writes**
 - Writes go to **memory** And WAL
 - **Sorted** memstore flushes to **disk**
 - **Sequential Reads**
 - From **memory**, **index**, **sorted** disk




predictable disk seeks



Data Model for Fast Writes, Reads

- Predictable disk lay out
- Minimize **disk seek**
- Get, Put by **row key**: **fast** access
- Scan by row **key range**: stored **sorted**, **efficient sequential** access for **key range**



Get key →

Scan start key,

Stop key →

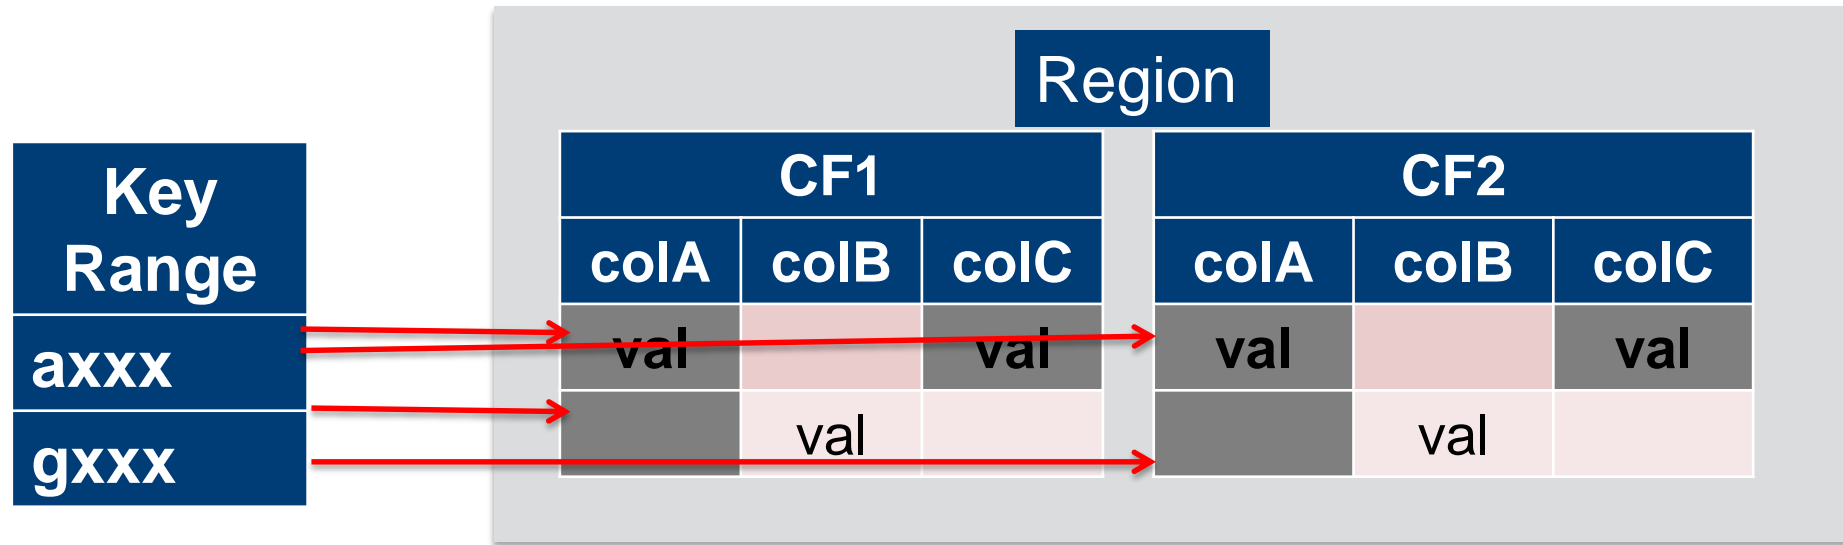
Region1 Key Range	
ra	
rx	

Region			
Key	CF1:Col	version	value
ra	cf1:ca	v1	1
rb	cf1:cb	v2	4
rb	cf1:cb	v1	3
rc	cf1:ca	v1	5

Region
Server



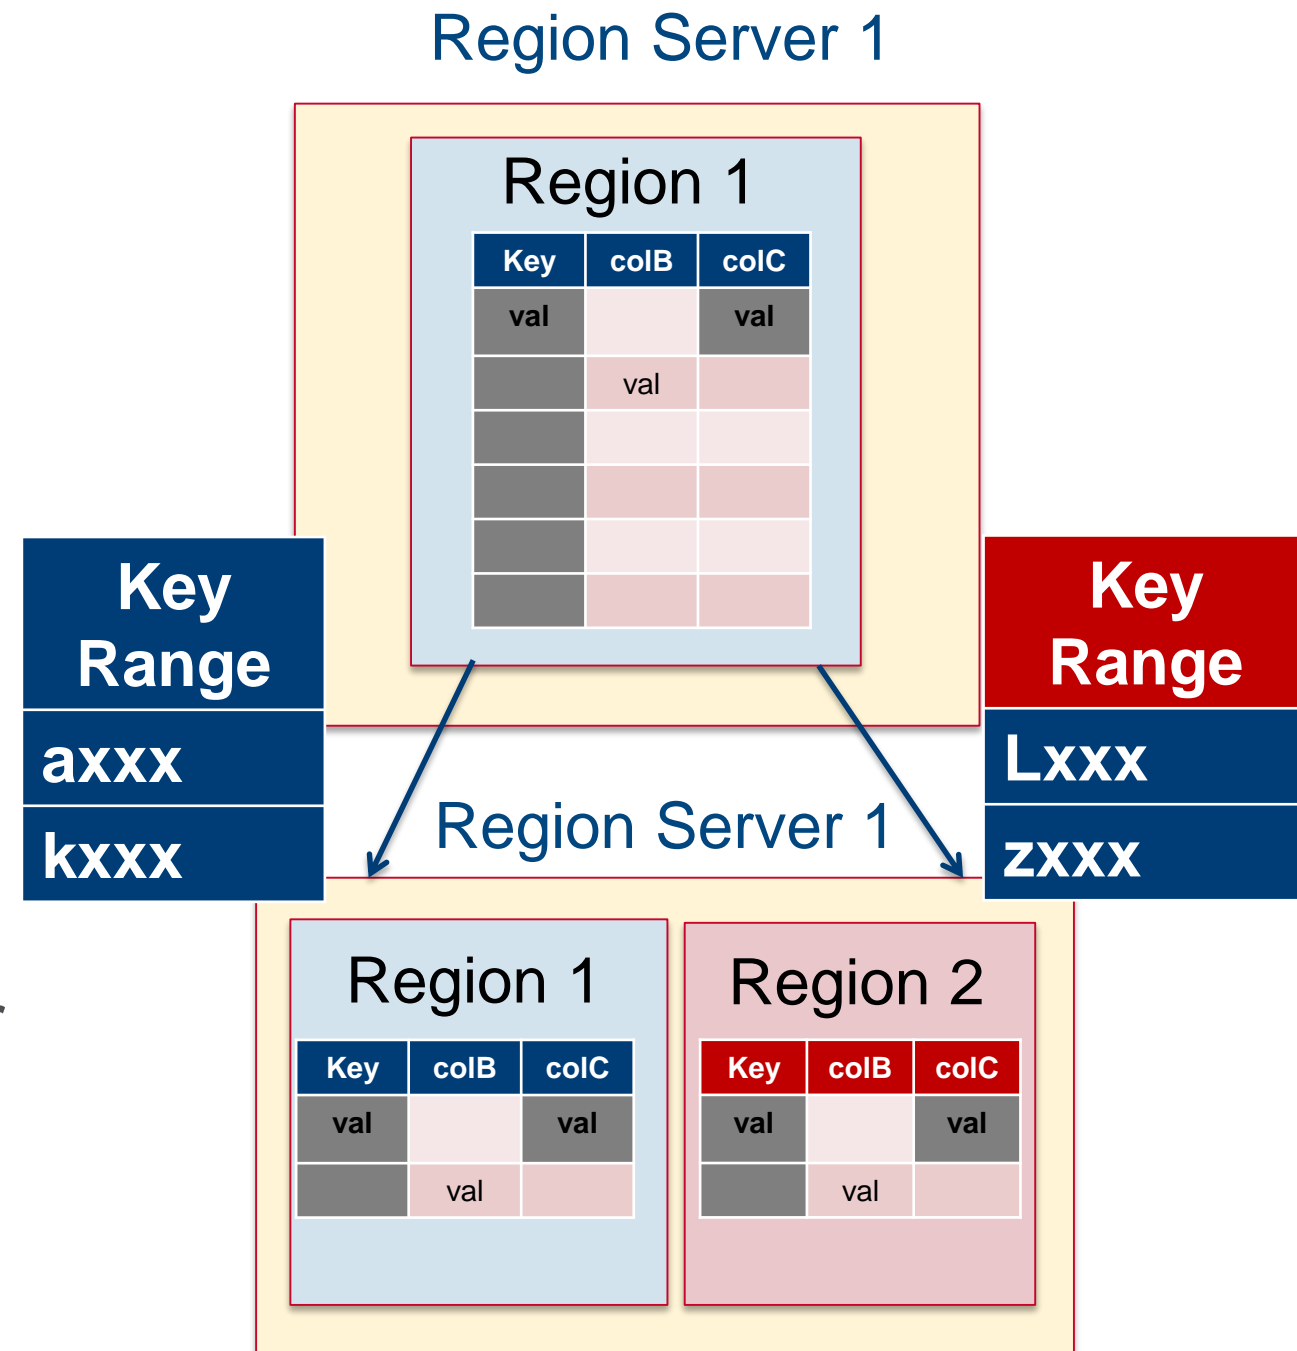
Region = contiguous keys



- **Regions** fundamental **partitioning/sharding** object.
- By **default**, on **table creation** **1 region** is created that holds the entire key range.
- When region becomes **too large**, **splits** into **two** child regions.
- Typical region size is a few GB, sometimes even 10G or 20G

Region Split

- The RegionServer splits a region
- daughter regions
 - each with $\frac{1}{2}$ of the regions **keys**.
 - opened in parallel on **same server**
- reports the split to the Master



HBase Use Cases



3 Main Use Case Categories

- Capturing Incremental data --Time Series Data
 - Hi Volume, Velocity **Writes**
- Information Exchange, Messaging
 - Hi Volume, Velocity **Write/Read**
- Content Serving, Web Application Backend
 - Hi Volume, Velocity **Reads**





3 Main Use Case Categories

- Time Series Data, Stuff with a Time Stamp
 - Sensor, System Metrics, Events, log files
 - Stock Ticker, User Activity
 - Hi Volume, Velocity Writes

OpenTSDB



Event time stamped
data

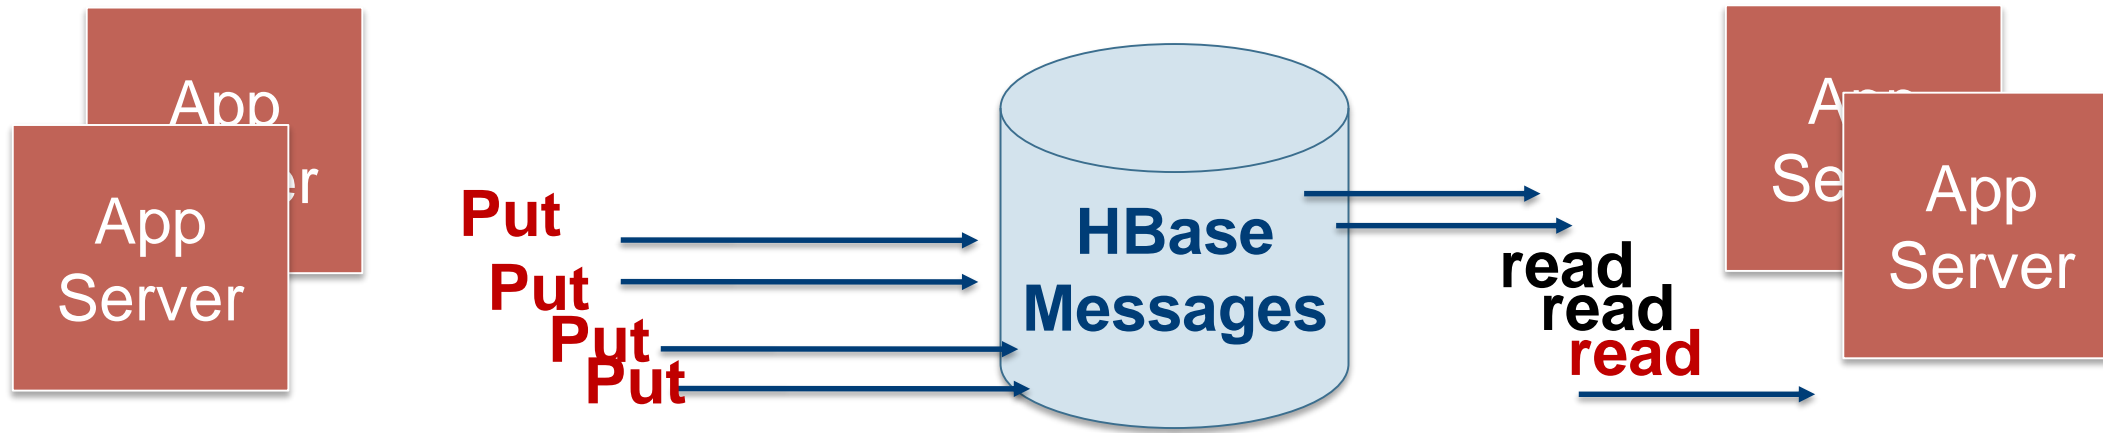


Data for real-time monitoring.



3 Main Use Case Categories

- Information Exchange
 - email, Chat, Inbox: **Facebook**
 - Hi Volume, Velocity Write/Read




3 Main Use Case Categories

- Content Serving, Web Application Backend
 - Online Catalog: Gap, World Library Catalog.
 - Search Index: ebay
 - Online Pre-Computed View: Groupon, Pinterest
 - Hi Volume, Velocity Reads

ebay

GROUPON

Bulk Import
Pre-Computed
Materialized View



Agenda

- Why do we need NoSQL / HBase?
- Overview of HBase & HBase data model
- HBase Architecture and data flow
- HBase Use Cases
- **Demo/Lab using HBase Shell**
 - Create tables and CRUD operations using MapR Sandbox
- HBase Java API to perform CRUD operations
 - Demo HBase Java API using Eclipse & MapR Sandbox



Prerequisite for Hands-On-Labs

Install MapR Sandbox

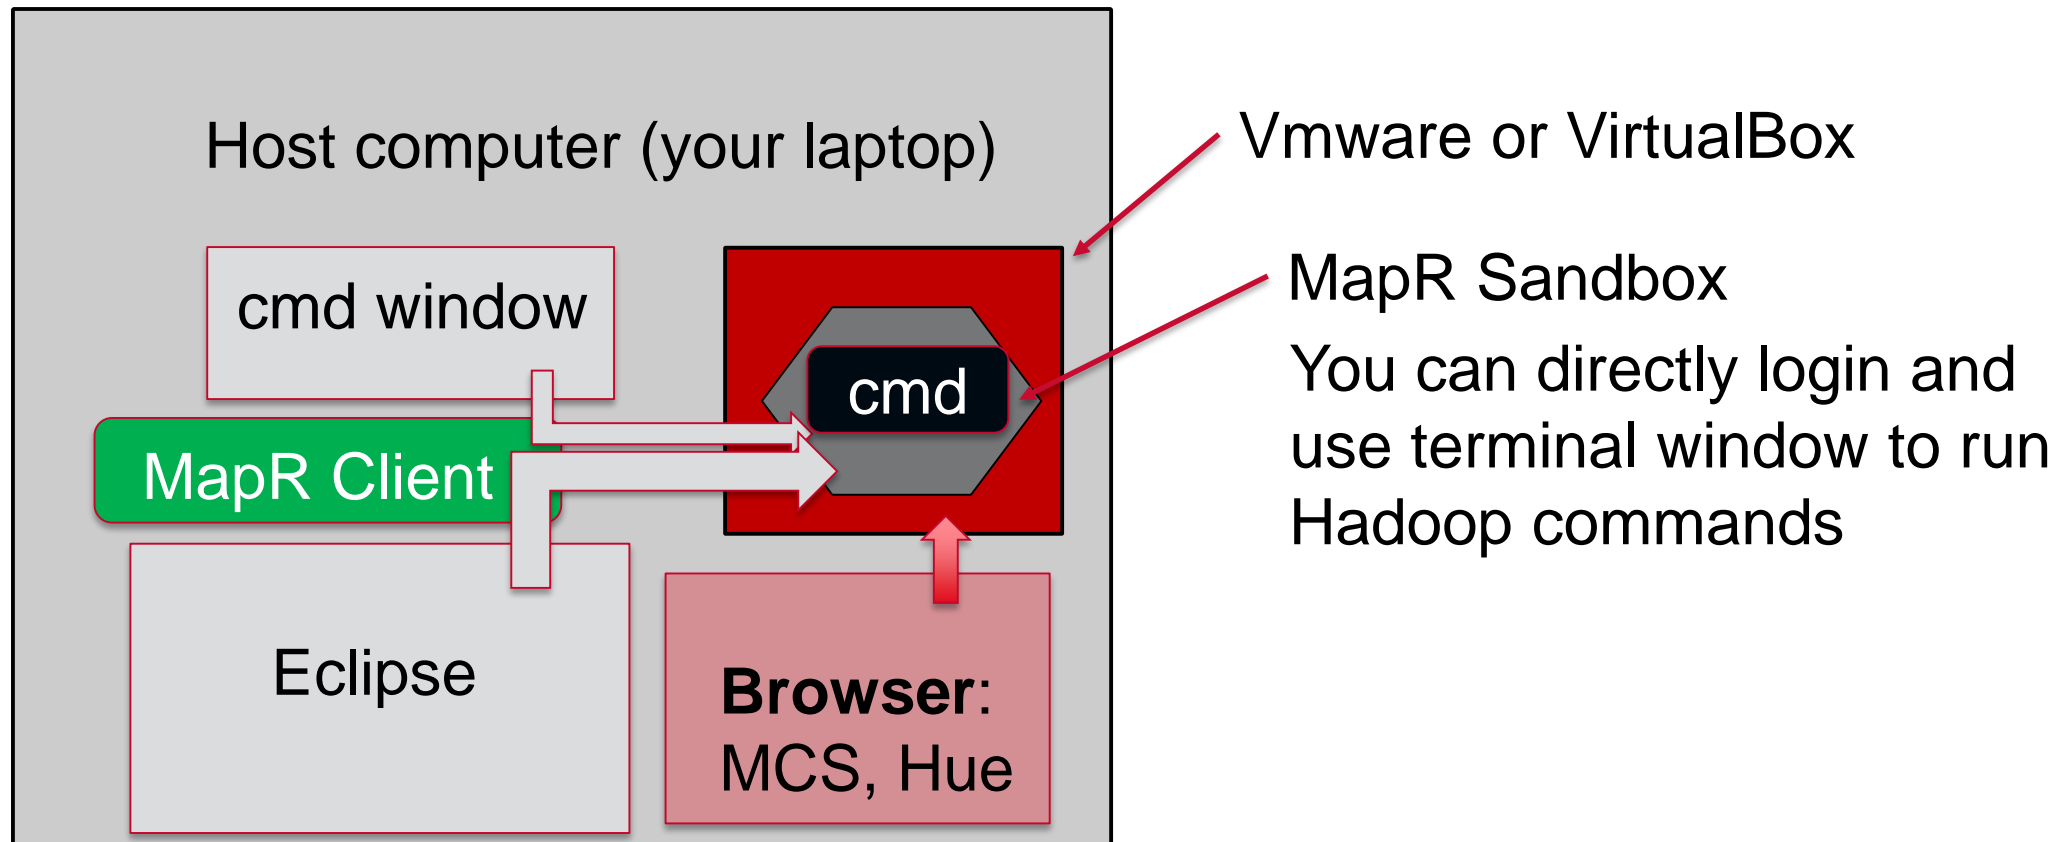
- Install a one-node MapR Sandbox on your laptop
- Install and configure Eclipse to develop HBase applications using Java API
- MapR Client is optional

Open Hbase_Tutorial_3_2015.pdf



What is MapR Sandbox and how to use it

- MapR Sandbox is a fully functional single-node Hadoop cluster running on a virtual machine



Lab Exercise

See [Lab_Hbase_Shell.pdf](#)

Start MapR Sandbox and log into the cluster

[user: mapr, passwd: mapr]

Use the HBase shell

>Hbase shell

hbase> *help*

hbase> *create '/user/mapr/mytable', {NAME =>'cf1'}*

hbase> *put '/user/mapr/mytable', 'row1', 'cf1:col1', 'datacf1c1v1'*

hbase> *get '/user/mapr/mytable', 'row1'*

hbase> *scan '/user/mapr/mytable'*

hbase> *describe '/user/mapr/mytable'*



Agenda

- Why do we need NoSQL / HBase?
- Overview of HBase & HBase data model
- HBase Architecture and data flow
- HBase Use Cases
- Demo/Lab using HBase Shell
 - Create tables and CRUD operations using MapR Sandbox
- **HBase Java API to perform CRUD operations**
 - Demo / Lab using Eclipse, HBase Java API & MapR Sandbox



HBase

Java API fundamentals to perform CRUD operations



Shoppingcart Application Requirements

- Need to create Tables: Shoppingcart & Inventory
- Perform **CRUD** operations on these tables
 - Create, Read, Update, and Delete items from these tables



Inventory & Shoppingcart Tables

Perform checkout operation for Mike

Inventory Table

	CF "stock "
	quantity
Pens	10
Notepads	21
Erasers	10
Pencils	40

Shoppingcart Table

	CF "items"		
	pens	notepads	erasers
Mike	1	2	3
John	3	4	5
Mary	1	2	5
Adam	5	4	0



Java API Fundamentals

- **CRUD operations**
 - **Get, Put, Delete, Scan, checkAndPut, checkAndDelete, Increment**
 - **KeyValue, Result, Scan – ResultScanner,**
 - **Batch Operations**



CRUD Operations Follow A Pattern (mostly)

- **common pattern**

- Instantiate object for an operation: `Put put = new Put(key)`
- Add attributes to specify **what to insert**: `put.add(...)`
- invoke operation with HTable: `myTable.put(put)`

```
// Insert value1 into rowKey in columnFamily:columnName1
Put put = new Put(rowKey);
put.add(columnFamily, columnName1, value1);
myTable.put(put);
```



Shopping Cart Table

Shoppingcart Table

	CF "items"		
	erasers	notepads	pens
Mike	3	2	1

Physical Storage

Key	CF:COL	ts	value
Mike	items:erasers	1391813876369	3
Mike	items:notepads	1391813876369	2
Mike	items:pens	1391813876369	1



Put Operation

Key	CF:COL	ts	value
Mike	items:erasers	1391813876369	3
Mike	items:notepads	1391813876369	2
Mike	items:pens	1391813876369	1

adding multiple column values to a row

```
byte [] tableName = Bytes.toBytes("/path/Shopping");
byte [] itemsCF = Bytes.toBytes("items");
byte [] penCol = Bytes.toBytes ("pens");
byte [] noteCol = Bytes.toBytes ("notes");
byte [] eraserCol = Bytes.toBytes ("erasers");
HTableInterface table = new HTable(hbaseConfig, tableName);

Put put = new Put("Mike");
put.add(itemsCF, penCol, Bytes.toBytes(1));
put.add(itemsCF, noteCol, Bytes.toBytes(2));
put.add(itemsCF, eraserCol, Bytes.toBytes(3));

table.put(put);
```



Get Example

Key	CF:COL	ts	value
Mike	items:erasers	1391813876369	3
Mike	items:notepads	1391813876369	2
Mike	items:pens	1391813876369	1

```
byte [] tableName = Bytes.toBytes("/user/user01/shoppingcart");
byte [] itemsCF = Bytes.toBytes("items");
byte [] penCol = Bytes.toBytes("pens");
HTableInterface table = new HTable(hbaseConfig, tableName);

Get get = new Get("Mike");

get.addColumn(itemsCF, penCol);

Result result = myTable.get(get);

byte[] val = result.getValue(itemsCF, penCol);

System.out.println("Value: " + Bytes.toLong(val));    //prints 1
```

Result Class

- A **Result** instance **wraps** data from a **row** returned from a **get** or a **scan** operation. Result wraps KeyValues

	Items:erasers	Items:notepads	Items:pens
Adam	0	4	5

- Result** toString() looks like this :

```
keyvalues={Adam/items:erasers/1391813876369/Put/vlen=8/ts=0,  
           Adam/items:notepads/1391813876369/Put/vlen=8/ts=0,  
           Adam/items:pens/1391813876369/Put/vlen=8/ts=0}
```

- The **Result** object provides methods to return **values**

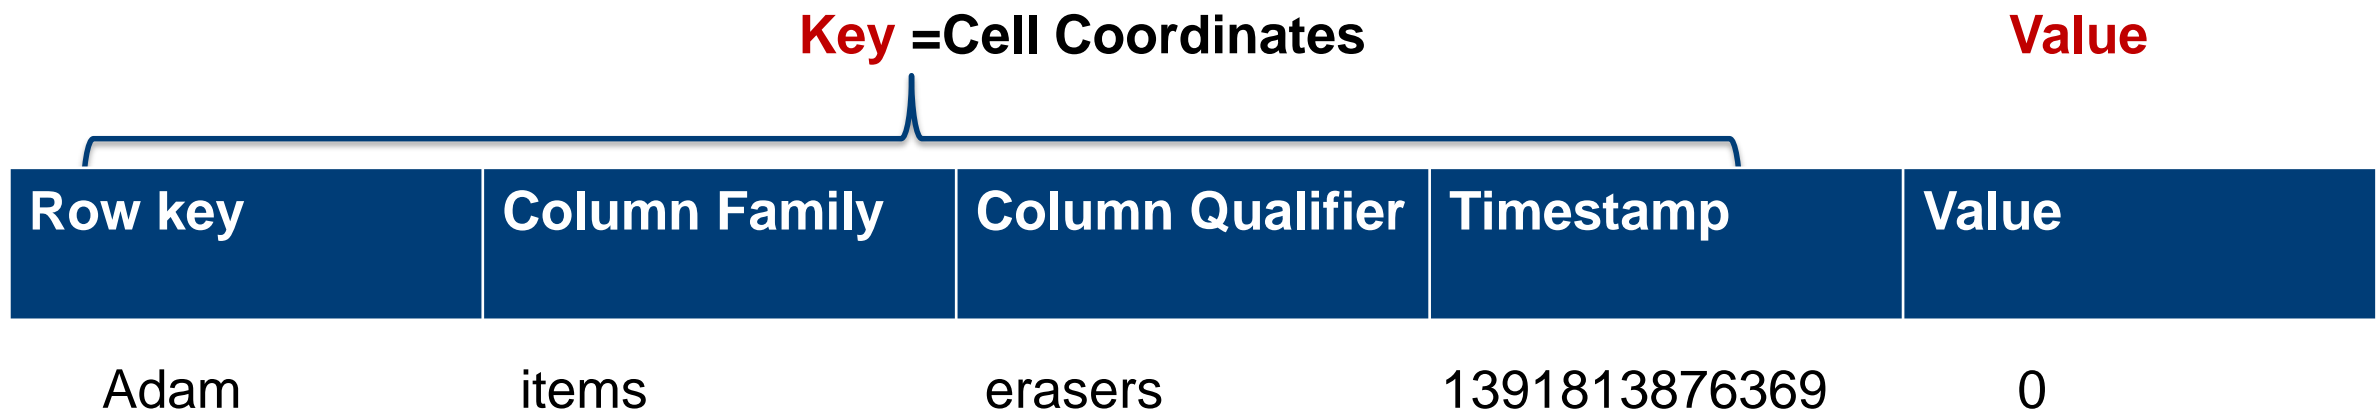
```
byte[] b = result.getValue(columnFamilyName, columnName1) ;
```

<http://hbase.apache.org/0.94/apidocs/org/apache/hadoop/hbase/client/Result.html>



KeyValue – The Fundamental HBase Type

- A **KeyValue** instance is a **cell** instance
 - Contains **Key** (cell coordinates) and the **Value** (data)
- Cell coordinates: Row key, Column family, Column qualifier, Timestamp
- **KeyValue** toString() looks like this :
Adam/items:erasers/1391813876369/Put/vlen=8/



Bytes class

<http://hbase.apache.org/0.94/apidocs/org/apache/hadoop/hbase/util/Bytes.html>


- org.apache.hadoop.hbase.util.Bytes
- Provides methods to convert Java types **to** and **from byte[]** arrays
- Support for
 - String, boolean, short, int, long, double, and float

```
byte[] bytesTable = Bytes.toBytes("Shopping");  
String table = Bytes.toString(bytesTable);  
  
byte[] amountBytes = Bytes.toBytes(10001);  
long amount = Bytes.toLong(amountBytes);
```



Scan Operation – Example

```
byte[] startRow=Bytes.toBytes("Adam");  
byte[] stopRow=Bytes.toBytes("N");  
  
Scan s = new Scan(startRow, stopRow);  
  
scan.addFamily(columnFamily);  
  
ResultScanner rs = myTable.getScanner(s);
```



ResultScanner - Example

Resultscanner provides **iterator-like** functionality

```
Scan scan = new Scan();
scan.addFamily(columnFamily);
ResultScanner scanner = myTable.getScanner(scan);
try {
    for (Result res : scanner) {
        System.out.println(res);
    }
} catch (Exception e) {
    System.out.println(e);
} finally {
    scanner.close();
}
```

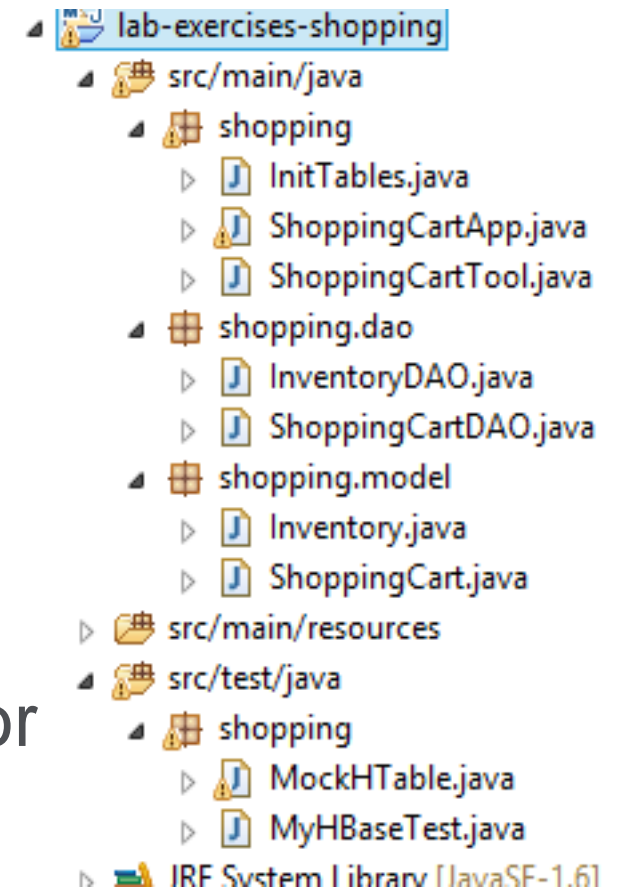
Calls scanner.next()

Always put in finally block



Lab Exercise Program Structure

- **ShoppingCartApp**– main class
- **InventoryDAO** – A DAO for the Inventory CRUD functionality
- **ShoppingCartDAO** – A DAO for the Inventory CRUD functionality
- **Inventory** – A Java object that holds data for a single Inventory row
- **ShoppingCart** – A Java object that holds data for a single Inventory row
- **MockHtable** – in memory test hbase table, allows to run code, debug without hbase running on a cluster or vm.



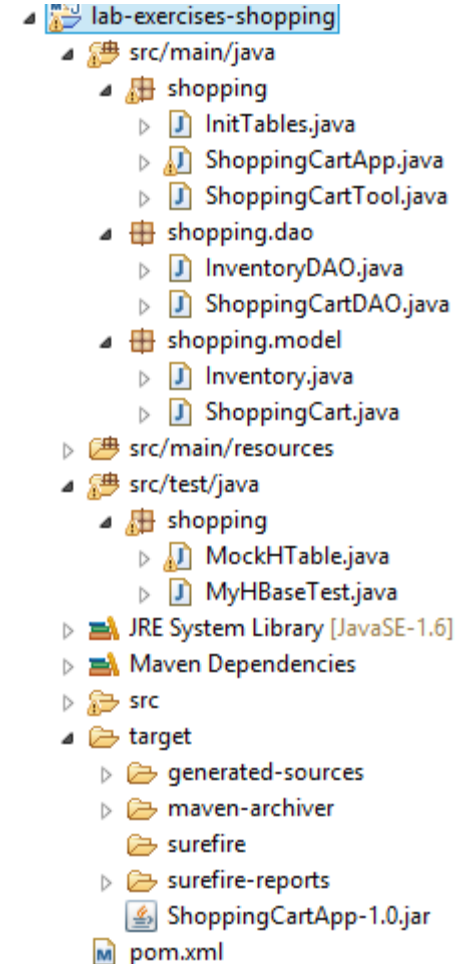
Lab Exercise

- See Lab_3_Java_API.pdf
 - Import the project “**lab-exercises-shopping**” into Eclipse
 - Setup creates Inventory and Shoppingcart Tables and inserts data
 - Use Get, Put, Scan, and Delete operations



Lab: Import, build

- **Download** the code
- **Import** Maven project lab-exercises-shopping into Eclipse
- **Build** : Run As -> Maven Install



Lab: run TestInventorySetup JUnit

- Select Test Class, Then Run As -> JUnit Test
- Uses MockHTable <https://gist.github.com/agaoglu/613217>

The screenshot displays an IDE interface with three main components:

- Project Tree (Left):** Shows a project named 'lab-exercise-shopping'. Under 'src/test/java', the 'shopping' package contains 'MockHTable.java' and 'TestInventorySetup.java'. 'TestInventorySetup.java' is selected.
- Context Menu (Center):** A right-click menu is open over 'TestInventorySetup.java'. The 'Run As' option at the bottom is expanded, showing 'JUnit 1 JUnit Test' as the selected action.
- Code Editor (Right):** Displays the content of 'TestInventorySetup.java'. The code includes imports for 'InventoryDAO' and 'MockHTable', and a test method 'testAddInventory' that uses 'assertEquals' to verify the results of 'addInventory'.

Summary

- Why do we need NoSQL / HBase?
- Overview of HBase & HBase data model
- HBase Architecture and data flow
- HBase Use Cases
- Demo/Lab using HBase Shell
 - Create tables and CRUD operations using MapR Sandbox
- HBase Java API to perform CRUD operations
 - Demo / Lab using Eclipse, HBase Java API & MapR Sandbox



References

- <http://hbase.apache.org/>
- <http://hbase.apache.org/0.94/apidocs/>
- <http://hbase.apache.org/0.94/apidocs/org/apache/hadoop/hbase/client/package-summary.html>
- <http://hbase.apache.org/book/book.html>

- <http://doc.mapr.com/display/MapR/MapR+Overview>
- <http://doc.mapr.com/display/MapR/M7+--+Native+Storage+for+MapR+Tables>
- <http://doc.mapr.com/display/MapR/MapR+Sandbox+for+Hadoop>
- <http://doc.mapr.com/display/MapR/Migrating+Between+Apache+HBase+Tables+and+MapR+Tables>




MapR-DB: Fast, Integrated NoSQL

Operational Analytics	Resilience/Business Continuity	Performance/Scale
Hadoop-enabled architecture <ul style="list-style-type: none">Operational/architectural simplicity, analytics on live dataHadoop-scale (petabytes)	Proven production readiness <ul style="list-style-type: none">Enterprise-grade HA/DRConsistent snapshotsIntegrated security	Consistent performance at any scale <ul style="list-style-type: none">High throughput (100M data points per second ingestion)Consistent low latency, no compaction delays, even at 95th and 99th percentiles (< 10ms in YCSB)



Mapr-DB Performance: Consistent, Low Latency



Q&A

Engage with us!

@mapr



maprtech

mapr-technologies



MapR

sreddy@mapr.com



maprtech

